

# Continual Zero-Shot Learning

Ivan Skorokhodov

March 17, 2020

# What is Continual Learning?

# What is Continual Learning?

- ▶ Modern neural networks are prone to *catastrophic forgetting*: they forget previous tasks while are learning new ones.

# What is Continual Learning?

- ▶ Modern neural networks are prone to *catastrophic forgetting*: they forget previous tasks while are learning new ones.
- ▶ Continual Learning tries to find ways to make model learn skills one by one in such a way that we do not forget previous skills

# What is Continual Learning?

- ▶ Modern neural networks are prone to *catastrophic forgetting*: they forget previous tasks while are learning new ones.
- ▶ Continual Learning tries to find ways to make model learn skills one by one in such a way that we do not forget previous skills
  - ▶ Example 1: a robot that travels the world and learn new skills. We want it not to forget previous skills while he is acquiring new ones.

# What is Continual Learning?

- ▶ Modern neural networks are prone to *catastrophic forgetting*: they forget previous tasks while are learning new ones.
- ▶ Continual Learning tries to find ways to make model learn skills one by one in such a way that we do not forget previous skills
  - ▶ Example 1: a robot that travels the world and learn new skills. We want it not to forget previous skills while he is acquiring new ones.
  - ▶ Example 2: a classification model is learning datasets one by one: we do not want its performance on previously learned datasets to decrease.

# Modern Continual Learning techniques

# Modern Continual Learning techniques

Modern CL techniques can be divided into three groups:



# Modern Continual Learning techniques

Modern CL techniques can be divided into three groups:

- ▶ Regularization-based ([4], [1], etc): detect the weights which are important for previous tasks and do not change them much in the future.

# Modern Continual Learning techniques

Modern CL techniques can be divided into three groups:

- ▶ Regularization-based ([4], [1], etc): detect the weights which are important for previous tasks and do not change them much in the future.
- ▶ Rehearsal-based ([2], [6], etc): store a part of previous data to replay it in the future.

# Modern Continual Learning techniques

Modern CL techniques can be divided into three groups:

- ▶ Regularization-based ([4], [1], etc): detect the weights which are important for previous tasks and do not change them much in the future.
- ▶ Rehearsal-based ([2], [6], etc): store a part of previous data to replay it in the future.
- ▶ Component-based ([5], [3], etc): divide your network into components, and let future tasks not to break components which are important for previous tasks.

# What is Zero-Shot Learning (ZSL)?

What data do we have:

# What is Zero-Shot Learning (ZSL)?

What data do we have:

- ▶ We will consider classification tasks from now on...

# What is Zero-Shot Learning (ZSL)?

What data do we have:

- ▶ We will consider classification tasks from now on...
- ▶ For each class  $c \in \mathcal{C}$  we are given an *attribute vector*  $a_c \in \mathcal{A}$  which describes the class:

# What is Zero-Shot Learning (ZSL)?

What data do we have:

- ▶ We will consider classification tasks from now on...
- ▶ For each class  $c \in \mathcal{C}$  we are given an *attribute vector*  $a_c \in \mathcal{A}$  which describes the class:
  - ▶ Imagine that we are classifying birds

# What is Zero-Shot Learning (ZSL)?

What data do we have:

- ▶ We will consider classification tasks from now on...
- ▶ For each class  $c \in \mathcal{C}$  we are given an *attribute vector*  $a_c \in \mathcal{A}$  which describes the class:
  - ▶ Imagine that we are classifying birds
  - ▶ Then for each bird  $a_c$  includes bird's characteristics: color of a tail, body size, length of a beak, etc



# What is Zero-Shot Learning (ZSL)?

What data do we have:

- ▶ We will consider classification tasks from now on...
- ▶ For each class  $c \in \mathcal{C}$  we are given an *attribute vector*  $a_c \in \mathcal{A}$  which describes the class:
  - ▶ Imagine that we are classifying birds
  - ▶ Then for each bird  $a_c$  includes bird's characteristics: color of a tail, body size, length of a beak, etc
- ▶ All classes are divided into *seen* and *unseen*:

# What is Zero-Shot Learning (ZSL)?

What data do we have:

- ▶ We will consider classification tasks from now on...
- ▶ For each class  $c \in \mathcal{C}$  we are given an *attribute vector*  $a_c \in \mathcal{A}$  which describes the class:
  - ▶ Imagine that we are classifying birds
  - ▶ Then for each bird  $a_c$  includes bird's characteristics: color of a tail, body size, length of a beak, etc
- ▶ All classes are divided into *seen* and *unseen*:
  - ▶ Seen dataset:  $D^s = \{X^s, Y^s, A^s\}$

# What is Zero-Shot Learning (ZSL)?

What data do we have:

- ▶ We will consider classification tasks from now on...
- ▶ For each class  $c \in \mathcal{C}$  we are given an *attribute vector*  $a_c \in \mathcal{A}$  which describes the class:
  - ▶ Imagine that we are classifying birds
  - ▶ Then for each bird  $a_c$  includes bird's characteristics: color of a tail, body size, length of a beak, etc
- ▶ All classes are divided into *seen* and *unseen*:
  - ▶ Seen dataset:  $D^s = \{X^s, Y^s, A^s\}$
  - ▶ Unseen dataset:  $D^u = \{X^u, Y^u, A^u\}$

# What is Zero-Shot Learning (ZSL)?

What data do we have:

- ▶ We will consider classification tasks from now on...
- ▶ For each class  $c \in \mathcal{C}$  we are given an *attribute vector*  $a_c \in \mathcal{A}$  which describes the class:
  - ▶ Imagine that we are classifying birds
  - ▶ Then for each bird  $a_c$  includes bird's characteristics: color of a tail, body size, length of a beak, etc
- ▶ All classes are divided into *seen* and *unseen*:
  - ▶ Seen dataset:  $D^s = \{X^s, Y^s, A^s\}$
  - ▶ Unseen dataset:  $D^u = \{X^u, Y^u, A^u\}$
- ▶ During training we have an access only to seen dataset  $D^s$ .

# What is Zero-Shot Learning (ZSL)?

What data do we have:

- ▶ We will consider classification tasks from now on...
- ▶ For each class  $c \in \mathcal{C}$  we are given an *attribute vector*  $a_c \in \mathcal{A}$  which describes the class:
  - ▶ Imagine that we are classifying birds
  - ▶ Then for each bird  $a_c$  includes bird's characteristics: color of a tail, body size, length of a beak, etc
- ▶ All classes are divided into *seen* and *unseen*:
  - ▶ Seen dataset:  $D^s = \{X^s, Y^s, A^s\}$
  - ▶ Unseen dataset:  $D^u = \{X^u, Y^u, A^u\}$
- ▶ During training we have an access only to seen dataset  $D^s$ .
  - ▶ Our goal is to learn to match images with class descriptions

# What is Zero-Shot Learning (ZSL)?

What data do we have:

- ▶ We will consider classification tasks from now on...
- ▶ For each class  $c \in \mathcal{C}$  we are given an *attribute vector*  $a_c \in \mathcal{A}$  which describes the class:
  - ▶ Imagine that we are classifying birds
  - ▶ Then for each bird  $a_c$  includes bird's characteristics: color of a tail, body size, length of a beak, etc
- ▶ All classes are divided into *seen* and *unseen*:
  - ▶ Seen dataset:  $D^s = \{X^s, Y^s, A^s\}$
  - ▶ Unseen dataset:  $D^u = \{X^u, Y^u, A^u\}$
- ▶ During training we have an access only to seen dataset  $D^s$ .
  - ▶ Our goal is to learn to match images with class descriptions
  - ▶ I.e. model learns to detect “blue tails”, “large heads”, “short beaks”, etc and not only concrete bird species

# What is Zero-Shot Learning (ZSL)?

What data do we have:

- ▶ We will consider classification tasks from now on...
- ▶ For each class  $c \in \mathcal{C}$  we are given an *attribute vector*  $a_c \in \mathcal{A}$  which describes the class:
  - ▶ Imagine that we are classifying birds
  - ▶ Then for each bird  $a_c$  includes bird's characteristics: color of a tail, body size, length of a beak, etc
- ▶ All classes are divided into *seen* and *unseen*:
  - ▶ Seen dataset:  $D^s = \{X^s, Y^s, A^s\}$
  - ▶ Unseen dataset:  $D^u = \{X^u, Y^u, A^u\}$
- ▶ During training we have an access only to seen dataset  $D^s$ .
  - ▶ Our goal is to learn to match images with class descriptions
  - ▶ I.e. model learns to detect "blue tails", "large heads", "short beaks", etc and not only concrete bird species
- ▶ At test time we evaluate model performance on unseen dataset  $D^u$

# What is Zero-Shot Learning (ZSL)?

What data do we have:

- ▶ We will consider classification tasks from now on...
- ▶ For each class  $c \in \mathcal{C}$  we are given an *attribute vector*  $a_c \in \mathcal{A}$  which describes the class:
  - ▶ Imagine that we are classifying birds
  - ▶ Then for each bird  $a_c$  includes bird's characteristics: color of a tail, body size, length of a beak, etc
- ▶ All classes are divided into *seen* and *unseen*:
  - ▶ Seen dataset:  $D^s = \{X^s, Y^s, A^s\}$
  - ▶ Unseen dataset:  $D^u = \{X^u, Y^u, A^u\}$
- ▶ During training we have an access only to seen dataset  $D^s$ .
  - ▶ Our goal is to learn to match images with class descriptions
  - ▶ I.e. model learns to detect "blue tails", "large heads", "short beaks", etc and not only concrete bird species
- ▶ At test time we evaluate model performance on unseen dataset  $D^u$
- ▶ Using the knowledge about how inputs and attributes correspond to each other we can detect birds that we have not seen before just based on their class description  $a_c$ .



## Modern ZSL techniques (for classification)

# Modern ZSL techniques (for classification)

- ▶ Embedding-based: build two embedder models:

# Modern ZSL techniques (for classification)

- ▶ Embedding-based: build two embedder models:
  - ▶ Model  $f(x)$  to embed images

# Modern ZSL techniques (for classification)

- ▶ Embedding-based: build two embedder models:
  - ▶ Model  $f(x)$  to embed images
  - ▶ Model  $h(a)$  to embed attributes

# Modern ZSL techniques (for classification)

- ▶ Embedding-based: build two embedder models:
  - ▶ Model  $f(x)$  to embed images
  - ▶ Model  $h(a)$  to embed attributes
  - ▶ Compute classification logits just by computing  $d(f(x), h(a_c))$  for each class  $c$  (here  $d$  is some distance function).

# Modern ZSL techniques (for classification)

- ▶ Embedding-based: build two embedder models:
  - ▶ Model  $f(x)$  to embed images
  - ▶ Model  $h(a)$  to embed attributes
  - ▶ Compute classification logits just by computing  $d(f(x), h(a_c))$  for each class  $c$  (here  $d$  is some distance function).
  - ▶ I.e. we just measure distance between an image embedding and attribute embeddings

# Modern ZSL techniques (for classification)

- ▶ Embedding-based: build two embedder models:
  - ▶ Model  $f(x)$  to embed images
  - ▶ Model  $h(a)$  to embed attributes
  - ▶ Compute classification logits just by computing  $d(f(x), h(a_c))$  for each class  $c$  (here  $d$  is some distance function).
  - ▶ I.e. we just measure distance between an image embedding and attribute embeddings
  - ▶ Challenges: how to embed images properly, how to embed attributes properly, what distance function to use, etc

# Modern ZSL techniques (for classification)

- ▶ Embedding-based: build two embedder models:
  - ▶ Model  $f(x)$  to embed images
  - ▶ Model  $h(a)$  to embed attributes
  - ▶ Compute classification logits just by computing  $d(f(x), h(a_c))$  for each class  $c$  (here  $d$  is some distance function).
  - ▶ I.e. we just measure distance between an image embedding and attribute embeddings
  - ▶ Challenges: how to embed images properly, how to embed attributes properly, what distance function to use, etc
- ▶ Generative-based



# Modern ZSL techniques (for classification)

- ▶ Embedding-based: build two embedder models:
  - ▶ Model  $f(x)$  to embed images
  - ▶ Model  $h(a)$  to embed attributes
  - ▶ Compute classification logits just by computing  $d(f(x), h(a_c))$  for each class  $c$  (here  $d$  is some distance function).
  - ▶ I.e. we just measure distance between an image embedding and attribute embeddings
  - ▶ Challenges: how to embed images properly, how to embed attributes properly, what distance function to use, etc
- ▶ Generative-based
  - ▶ Train a conditional generative model that will learn to generate images based on class descriptions

# Modern ZSL techniques (for classification)

- ▶ Embedding-based: build two embedder models:
  - ▶ Model  $f(x)$  to embed images
  - ▶ Model  $h(a)$  to embed attributes
  - ▶ Compute classification logits just by computing  $d(f(x), h(a_c))$  for each class  $c$  (here  $d$  is some distance function).
  - ▶ I.e. we just measure distance between an image embedding and attribute embeddings
  - ▶ Challenges: how to embed images properly, how to embed attributes properly, what distance function to use, etc
- ▶ Generative-based
  - ▶ Train a conditional generative model that will learn to generate images based on class descriptions
  - ▶ I.e. GAN model that can generate a pigeon given description "grey feather, white head, short legs, etc"

# Modern ZSL techniques (for classification)

- ▶ Embedding-based: build two embedder models:
  - ▶ Model  $f(x)$  to embed images
  - ▶ Model  $h(a)$  to embed attributes
  - ▶ Compute classification logits just by computing  $d(f(x), h(a_c))$  for each class  $c$  (here  $d$  is some distance function).
  - ▶ I.e. we just measure distance between an image embedding and attribute embeddings
  - ▶ Challenges: how to embed images properly, how to embed attributes properly, what distance function to use, etc
- ▶ Generative-based
  - ▶ Train a conditional generative model that will learn to generate images based on class descriptions
  - ▶ I.e. GAN model that can generate a pigeon given description "grey feather, white head, short legs, etc"
  - ▶ At test time generate a lot of synthetic images, then train a classifier based on this synthetic dataset

# Modern ZSL techniques (for classification)

- ▶ Embedding-based: build two embedder models:
  - ▶ Model  $f(x)$  to embed images
  - ▶ Model  $h(a)$  to embed attributes
  - ▶ Compute classification logits just by computing  $d(f(x), h(a_c))$  for each class  $c$  (here  $d$  is some distance function).
  - ▶ I.e. we just measure distance between an image embedding and attribute embeddings
  - ▶ Challenges: how to embed images properly, how to embed attributes properly, what distance function to use, etc
- ▶ Generative-based
  - ▶ Train a conditional generative model that will learn to generate images based on class descriptions
  - ▶ I.e. GAN model that can generate a pigeon given description "grey feather, white head, short legs, etc"
  - ▶ At test time generate a lot of synthetic images, then train a classifier based on this synthetic dataset
  - ▶ Challenges: how to train a good conditional generative model?

# Modern ZSL techniques (for classification)

- ▶ Embedding-based: build two embedder models:
  - ▶ Model  $f(x)$  to embed images
  - ▶ Model  $h(a)$  to embed attributes
  - ▶ Compute classification logits just by computing  $d(f(x), h(a_c))$  for each class  $c$  (here  $d$  is some distance function).
  - ▶ I.e. we just measure distance between an image embedding and attribute embeddings
  - ▶ Challenges: how to embed images properly, how to embed attributes properly, what distance function to use, etc
- ▶ Generative-based
  - ▶ Train a conditional generative model that will learn to generate images based on class descriptions
  - ▶ I.e. GAN model that can generate a pigeon given description "grey feather, white head, short legs, etc"
  - ▶ At test time generate a lot of synthetic images, then train a classifier based on this synthetic dataset
  - ▶ Challenges: how to train a good conditional generative model?
  - ▶ Currently performs better than embedding-based approaches

# Continual Zero-Shot Learning

# Continual Zero-Shot Learning

- ▶ Project: let's use class attributes to improve the performance on future tasks (and this also should improve past performance)

# Continual Zero-Shot Learning

- ▶ Project: let's use class attributes to improve the performance on future tasks (and this also should improve past performance)
- ▶ Why?



# Continual Zero-Shot Learning

- ▶ Project: let's use class attributes to improve the performance on future tasks (and this also should improve past performance)
- ▶ Why?
  - ▶ A robot should be able to understand things it has never seen before but only heard about.

# Continual Zero-Shot Learning

- ▶ Project: let's use class attributes to improve the performance on future tasks (and this also should improve past performance)
- ▶ Why?
  - ▶ A robot should be able to understand things it has never seen before but only heard about.
  - ▶ And it shouldn't forget previously seen objects while doing so...

# Continual Zero-Shot Learning

- ▶ Project: let's use class attributes to improve the performance on future tasks (and this also should improve past performance)
- ▶ Why?
  - ▶ A robot should be able to understand things it has never seen before but only heard about.
  - ▶ And it shouldn't forget previously seen objects while doing so...
  - ▶ And the set of attribute descriptions can grow over time...

# Continual Zero-Shot Learning

- ▶ Project: let's use class attributes to improve the performance on future tasks (and this also should improve past performance)
- ▶ Why?
  - ▶ A robot should be able to understand things it has never seen before but only heard about.
  - ▶ And it shouldn't forget previously seen objects while doing so...
  - ▶ And the set of attribute descriptions can grow over time...
  - ▶ In some sense, it is "the next" step of ZSL

# Continual Zero-Shot Learning

- ▶ Project: let's use class attributes to improve the performance on future tasks (and this also should improve past performance)
- ▶ Why?
  - ▶ A robot should be able to understand things it has never seen before but only heard about.
  - ▶ And it shouldn't forget previously seen objects while doing so...
  - ▶ And the set of attribute descriptions can grow over time...
  - ▶ In some sense, it is "the next" step of ZSL
- ▶ I.e. build a *zero-shot* model that is trained in a *continual learning* fashion

What we have already tried

## What we have already tried

- ▶ Attempt #1: use some simple (but working) generative-based ZSL approach

# What we have already tried

- ▶ Attempt #1: use some simple (but working) generative-based ZSL approach
  - ▶ Problem: generative models are hard to train in CL scenario (it is slow and they forget things fast)



# What we have already tried

- ▶ Attempt #1: use some simple (but working) generative-based ZSL approach
  - ▶ Problem: generative models are hard to train in CL scenario (it is slow and they forget things fast)
- ▶ Attempt #2: use some simple (but working) embeddings-based ZSL approach







# What we have already tried

- ▶ Attempt #1: use some simple (but working) generative-based ZSL approach
  - ▶ Problem: generative models are hard to train in CL scenario (it is slow and they forget things fast)
- ▶ Attempt #2: use some simple (but working) embeddings-based ZSL approach
  - ▶ Problem #1: It was already explored in previous literature (in A-GEM paper)

# What we have already tried

- ▶ Attempt #1: use some simple (but working) generative-based ZSL approach
  - ▶ Problem: generative models are hard to train in CL scenario (it is slow and they forget things fast)
- ▶ Attempt #2: use some simple (but working) embeddings-based ZSL approach
  - ▶ Problem #1: It was already explored in previous literature (in A-GEM paper)
  - ▶ Problem #2: Attributes do not have optimal “discriminative” power

# References

-  Rahaf Aljundi et al. "Memory Aware Synapses: Learning what (not) to forget". In: *CoRR* abs/1711.09601 (2017).
-  Arslan Chaudhry et al. "Efficient Lifelong Learning with A-GEM". In: *International Conference on Learning Representations*. 2019.
-  Chrisantha Fernando et al. "PathNet: Evolution Channels Gradient Descent in Super Neural Networks". In: *CoRR* abs/1701.08734 (2017).
-  James Kirkpatrick et al. "Overcoming catastrophic forgetting in neural networks". In: *Proceedings of the National Academy of Sciences* 114.13 (2017), pp. 3521–3526.
-  Joan Serra et al. "Overcoming Catastrophic Forgetting with Hard Attention to the Task". In: *ICML*. Vol. 80. *Proceedings of Machine Learning Research*. PMLR, Oct. 2018, pp. 4548–4557.
-  Chenshen Wu et al. "Memory Replay GANs: Learning to Generate New Categories without Forgetting". In: *Advances in Neural Information Processing Systems 31*. Curran Associates, Inc., 2018, pp. 5962–5972.