

# Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains<sup>1</sup>

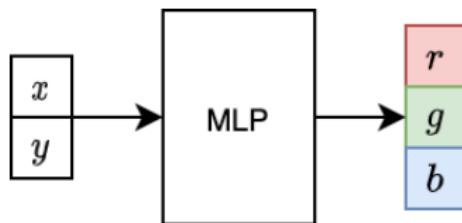
July 8, 2020

---

<sup>1</sup>*Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains* by Tancik et al., 2020

# Overview if INRs

- ▶ Recently, there appeared a new paradigm of representing signals
- ▶ One can represent images/3D-objects/video/audio/etc as *implicit neural representations* (INR)
- ▶ INR is a neural network  $\Phi(v)$  which takes a coordinates vector  $v$  and produces a pixel/voxel/frequency/etc value at that point:



**Figure:** INR for 2D images: at each coordinate we predict a pixel value

- ▶ I.e. one neural network corresponds to a single image.

## Overview of the paper

- ▶ Authors showed that standard MLPs fail to learn high frequency functions (both in theory and in practice)
- ▶ They showed that using “fourier features” overcomes the issue
- ▶ They test their ideas on a range of INR tasks: reconstructing images/3D-shapes/MRI-images/etc
- ▶ They used JAX to implement this all

## Intuition. Part 1/3: the foreword

- ▶ Let's roughly divide images into low-frequency images and high-frequency images:
  - ▶ In a low-frequency image, things change slow from pixel to pixel
  - ▶ In a high-frequency image, things change fast



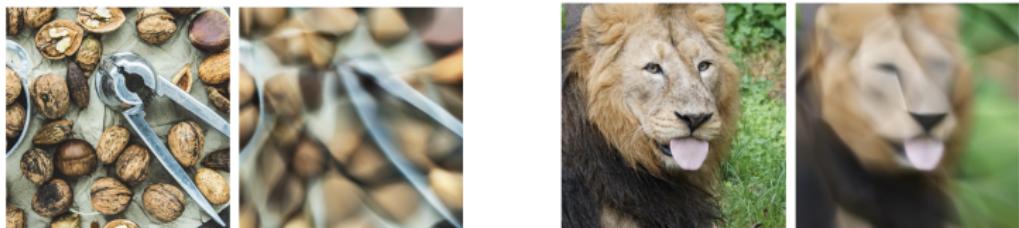
(a) A low-frequency image



(b) A high-frequency image

## Intuition. Part 2/3: the problem

- ▶ Neural networks are bad at distinguishing small values between each other
- ▶ For example, if your input is 0.03 and you replace it with 0.05, your model will likely not notice any difference
- ▶ And this is the main cause of why MLPs fail to model high-frequency images:
  - ▶ Our input coordinates for an INR lie in  $[0, 1]$  domain
  - ▶ For high-frequency images, output changes a lot when an input changes slightly
  - ▶ Our model is incapable to detect small changes in the input
  - ▶ The result: MLPs "oversmooth" high-frequency images



**Figure:** INR is trained to reconstruct an image: MLP oversmoothes high-frequency images

## Intuition. Part 3/3: the solution

- ▶ How can we help the model to detect small changes in an input?
  - ▶ Step 1: multiply by a very large number
  - ▶ Step 2: apply  $\sin(x)$  or  $\cos(x)$

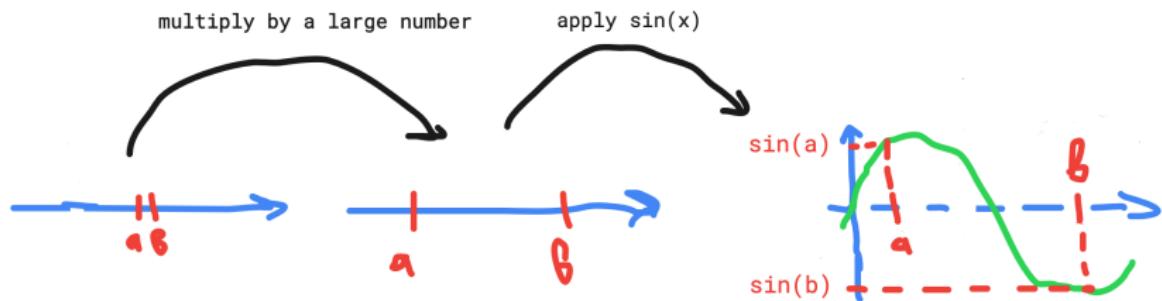


Figure: Intuition of the positional encoding trick

- ▶ Why do we need sines/cosines?
  - ▶ Neural networks do not like large-magnitude values (due to exploding variance)
  - ▶ Authors also show that this allows a model to have a shift-invariance property

## Kernel regression

- ▶ Kernel regression is the following machine learning algorithm:
  - ▶ Obtain a training dataset  $\mathbf{X}, \mathbf{y}$
  - ▶ Obtain a kernel function  $k(\mathbf{x}, \mathbf{x}')$  (it's like a similarity function)
  - ▶ Compute kernel matrix  $\mathbf{K}$  where  $K_{ij} = k(\mathbf{x}_i, \mathbf{x}_j)$
  - ▶ Compute the value for a test point  $\mathbf{x}$  by:

$$\hat{f}(\mathbf{x}) = \sum_{i=1}^n \left( \mathbf{K}^{-1} \mathbf{y} \right)_i k(\mathbf{x}_i, \mathbf{x}) \quad (1)$$

## NTK theory

NTK (Neural Tangent Kernel) theory explores the behaviour of infinite-width MLPs:

- ▶ If the learning rate is small then the MLP simplifies to kernel regression with:

$$k_{\text{NTK}}(\mathbf{x}_i, \mathbf{x}_j) = \mathbb{E}_{\theta \sim \mathcal{N}} \left\langle \frac{\partial f(\mathbf{x}_i; \theta)}{\partial \theta}, \frac{\partial f(\mathbf{x}_j; \theta)}{\partial \theta} \right\rangle$$

- ▶ Matrix  $\mathbf{K}$  is poorly conditioned for a regular MLP, i.e. it has few large eigenvalues and others are small
- ▶ NTK shows that the components of the target function that correspond to large eigenvalues are learned much faster
- ▶ Components with large eigenvalues have lower frequency
- ▶ This makes the learning of high-frequency patterns too slow
- ▶ Authors propose to preprocess input coordinates with Fourier mapping:
  - ▶ This would make  $k_{\text{NTK}}$  be better conditioned
  - ▶ This would make  $k_{\text{NTK}}$  stationary<sup>2</sup> which gives invariance to shifts

---

<sup>2</sup>i.e.  $k(\mathbf{x}, \mathbf{x}') = h(\mathbf{x} - \mathbf{x}')$  for some  $h$

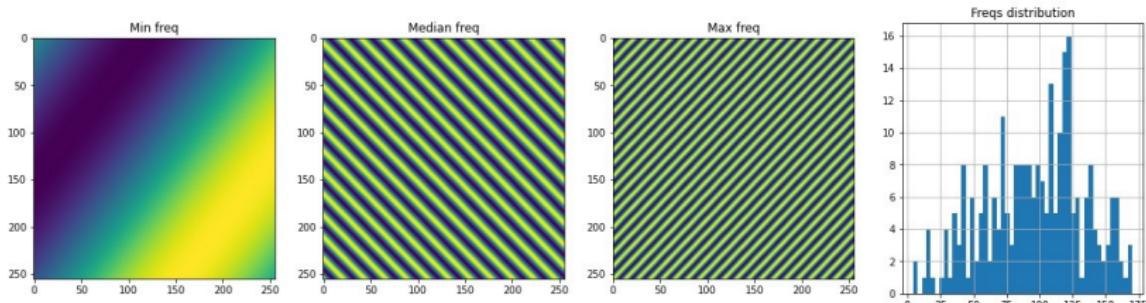
# Fourier mapping for INRs

- ▶ Authors propose to embed the coordinates with:

$$\gamma(\mathbf{v}) = [\cos(2\pi \mathbf{b}_1^T \mathbf{v}), \sin(2\pi \mathbf{b}_1^T \mathbf{v}), \dots, \cos(2\pi \mathbf{b}_m^T \mathbf{v}), \sin(2\pi \mathbf{b}_m^T \mathbf{v})]^T$$

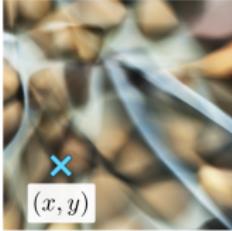
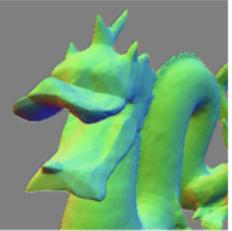
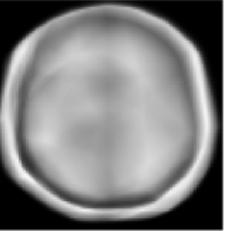
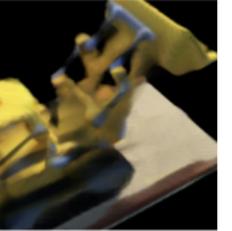
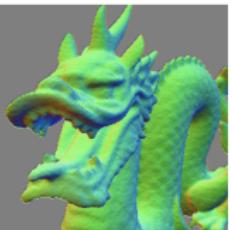
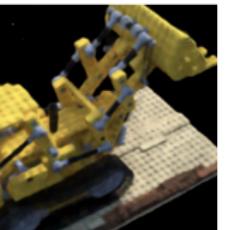
where  $\mathbf{b}_i$  are random vectors of size  $d$  (for images,  $d = 2$ )

- ▶ This makes the NTK kernel stationary and have a better spectrum
- ▶ Authors sample  $\mathbf{b}_i$  from  $\mathcal{N}(0, \sigma^2)$  where  $\sigma^2$  is a large number



**Figure:** Visualization of the embeddings for some of the vectors  $\mathbf{b}_i$  for  $256 \times 256$  coordinates. Vectors  $\mathbf{b}_i$  are sampled from uniform distribution.

# Results

No Fourier features $\gamma(\mathbf{v}) = \mathbf{v}$	 A photograph of walnuts and a nutcracker. A blue cross indicates a 2D point $(x, y)$ on the walnuts.	 A 3D surface reconstruction of a dragon head, colored with a gradient from blue to green.	 A grayscale MRI slice showing a brain cross-section.	 A 3D reconstruction of a yellow toy truck on a checkered surface.
With Fourier features $\gamma(\mathbf{v}) = \text{FF}(\mathbf{v})$	 A photograph of walnuts and a nutcracker, showing more detail and texture than the no-features version.	 A 3D surface reconstruction of a dragon head, showing sharper edges and more detail than the no-features version.	 A grayscale MRI slice showing a brain cross-section, with more fine-grained detail compared to the no-features version.	 A 3D reconstruction of a yellow toy truck on a checkered surface, showing a smoother and more accurate fit to the original image than the no-features version.

## Conclusion

- ▶ Authors provide an interesting (but not very transparent) theoretical analysis of INRs
- ▶ Their theoretical claims are confirmed by empirical evidence
- ▶ They achieve strong results on several domains
- ▶ One needs to search for optimal  $\sigma$  value when sampling  $\boldsymbol{b}$  from  $\mathcal{N}(0, \sigma^2)$