

# Generative Memory for Continual Learning

Ivan Skorokhodov

December 1, 2019

# Contents

1. Deep Generative Replay
2. MeRGAN
3. Dynamic Generative Memory
4. Latent Generative Memory

# Deep Generative Replay (DGR) <sup>1</sup>

Main idea (1/3)

- ▶ Train a generator  $G_1$ , train a classifier  $C_1$  for task #1
- ▶ For task  $t > 1$  generate images with  $G_{t-1}$ , generate labels with  $C_{t-1}$  to obtain a dataset  $\hat{D}_{:t}$
- ▶ Train on both  $\hat{D}_{:t}$  and  $D_t$  (real data for task  $t$ ) jointly
- ▶ Note: it's not clear from the paper if they trained  $C_t$  on the logits of  $C_{t-1}$  or its one-hot predictions
- ▶ Note: it's a bit odd that they do not train conditional generator


$$\mathcal{L} = r \mathbb{E}_{(\mathbf{x}, \mathbf{y}) \sim D_t} [L(C_t(\mathbf{x}), \mathbf{y})] + (1 - r) \mathbb{E}_{\mathbf{x}' \sim G_{t-1}} [L(C_t(\mathbf{x}'), C_{t-1}(\mathbf{x}'))]$$

---

<sup>1</sup> “Continual Learning with Deep Generative Replay” by Shin et al., NeurIPS 2017


# Deep Generative Replay (DGR)

## Illustration (2/3)



# Deep Generative Replay (DGR)

Results on permuted MNIST (3/3)



- ▶ 5 tasks
- ▶ ER — Joint Multi-Task baseline
- ▶ Noise — feeding random noise instead of images into  $C_{t-1}$  to distill knowledge

# Memory Replay GAN <sup>2</sup>

Main idea (1/3)

- ▶ Idea is simple: train a generative memory  $G_t$ , save its snapshot before each new task and distill its knowledge into a new one  $G_{t+1}$
- ▶ There are two ways to distill the knowledge
  - ▶ Generate synthetic data and mix it into a new one  $S'_t$  (Joint Retraining)
  - ▶ Perform real knowledge distillation (Replay Alignment):

$$\mathcal{L}_G = L_G(\theta_t, S_t) + \lambda \mathbb{E}_{z \sim p_z, c \sim U(0, t-1)} \left[ \|G_t(z, c) - G_{t-1}(z, c)\|^2 \right]$$

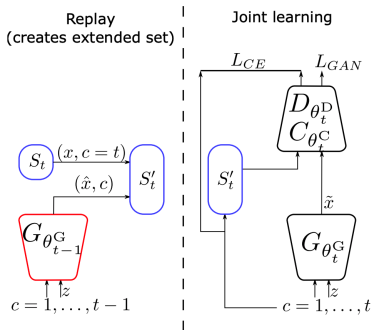
- ▶ Also train a classifier on top of GAN

---

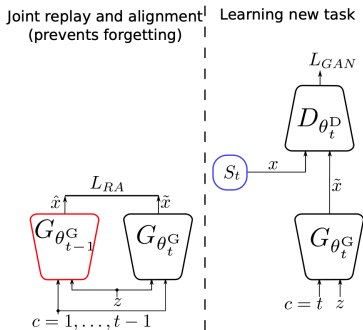
<sup>2</sup> “Memory Replay GANs: Learning to Generate New Categories without Forgetting” by Wu et al., NeurIPS 2018

# Memory Replay GAN (MeRGAN)

## Illustration (2/3)




(a) Joint retraining with replay



(b) Replay alignment

# Memory Replay GAN (MeRGAN)

Results (3/3)



- SFT (sequential fine tuning) is no tricks at all
- Replay Alignment tends to work better
- Authors are not clear about how they have measured the accuracy, as far as I got — they have trained a classifier on real data and measured its performance on the fake data.



# Dynamic Generative Memory (DGM) <sup>4</sup>

## Main idea (1/3)

- ▶ Authors consider class-incremental learning (just as we do): classes arrive sequentially and we evaluate the performance on all the tasks
- ▶ They do not run any kind of knowledge distillation and follow HAT<sup>3</sup> approach instead
- ▶ More precisely, for task  $t$ , for layer  $l$  of the generator they train a binary mask  $m_l^t$  and multiply layer's weights on this mask
- ▶ Binary mask  $m_l^t$  is regularized to be sparse
- ▶ Previously learned weights are not updated in the future (but network can learn to ignore them by learning the corresponding mask)
- ▶ But compared to HAT, authors are cheating: they add new neurons to the generator after each task to preserve its capacity


---

<sup>3</sup>“Overcoming Catastrophic Forgetting with Hard Attention to the Task” by Serra et al., ICML 2018

<sup>4</sup>“Learning to Remember: A Synaptic Plasticity Driven Framework for Continual Learning” by Ostapenko et al., arxiv

# Dynamic Generative Memory (DGM)

## Illustration (2/3)



# Dynamic Generative Memory (DGM)

## Results (3/3)

	Method	MNIST (%)		SVHN(%)		CIFAR10(%)		ImageNet-50(%)	
		$A_5$	$A_{10}$	$A_5$	$A_{10}$	$A_5$	$A_{10}$	$A_{30}$	$A_{50}$
Episodic memory	JT	99.87	99.24	92.99	88.72	83.40	77.82	57.35	49.88
	iCarl-S [22]	-	55.8	-	-	-	-	29.38	<b>28.98</b>
	EWC-S[9]	-	79.7	-	-	-	-	-	-
	RWalk-S[2]	-	82.5	-	-	-	-	-	-
	PI-S [34]	-	78.7	-	-	-	-	-	-
Generat. memory	EWC-M [28]	70.62	77.03	39.84	33.02	-	-	-	-
	DGR [30]	90.39	85.40	61.29	47.28	-	-	-	-
	MeRGAN [31]	98.19	97.00	80.90	66.78	-	-	-	-
	DGMw (ours)	98.75	96.46	<b>83.93</b>	<b>74.38</b>	<b>72.45</b>	<b>56.21</b>	<b>32.14</b>	17.82
	DGMa (ours)	<b>99.17</b>	<b>97.92</b>	81.07	66.89	71.91	51.75	25.93	15.16

- ▶ Here  $A_n$  is the performance on  $n$  previously seen classes
- ▶ For ImageNet-50 they train for 5 tasks, 10 classes per task
- ▶ They do not state it clearly, but as far as I got they use 5 and 10 tasks for other datasets


# Latent Generative Memory

## Main idea (1/2)

- ▶ Let's get rid of separate knowledge distillation step (it is questionable both biologically and practically)
- ▶ So let's train GM and Classifier jointly
- ▶ Since training GM in the visual space is tough, let's train it in the feature space
- ▶ Make the GM reside in “deep” layers of the Classifier and hallucinate

# Latent Generative Memory

## Illustration (2/2)



- ▶ For task  $t = 1$  we train the model normally
- ▶ For task  $t > 1$  generate a lot of fake memories with  $G_{t-1}$  of previously seen classes
- ▶ Train Classifier to correctly distinguish these fake memories
- ▶ Question #1: how to avoid knowledge distillation for Generator?
- ▶ Question #2: what if Encoder will start changing the embedding manifold? Then our fake memories will not correspond to actual embeddings. Maybe we can introduce prototypes to resolve this?

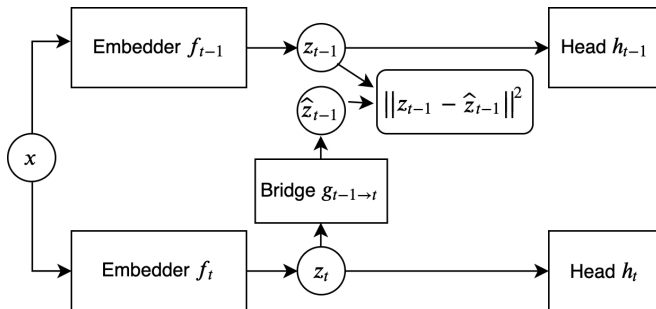
# Latent Space Alignment

## Main idea (1/2)

- ▶ It's not about generative memory (but can be useful)
- ▶ Imagine our classifier is  $h(f(x))$ , where  $f(x)$  is an embedder and  $h(z)$  is a head
- ▶ We work in multi-headed setup, i.e. we have a separate head  $h_t(z)$  for each task  $t$
- ▶ In this setup forgetting occurs when embedder  $f(x)$  changes
- ▶ Imagine that head  $h_{t-1}(z)$  was operating in embedding space  $Z_{t-1}$ , but task  $t$  changed it and  $h_t(z)$  operates in  $Z_t$
- ▶ Let's train a "bridge" function  $g_{t \rightarrow t-1} : Z_t \rightarrow Z_{t-1}$  which will convert  $Z_t$  to  $Z_{t-1}$
- ▶ Having good bridges between all the latent spaces  $Z_T \rightarrow Z_{T-1} \rightarrow \dots \rightarrow Z_2 \rightarrow Z_1$  we'll be able to compute predictions with embedding  $f_T(x)$  without forgetting
- ▶ An interesting consequence is that we can train 10 models on 10 tasks in parallel and then just align their latent spaces

# Latent Space Alignment

## Illustration (2/2)



- After we have trained a bridge  $g_{t \rightarrow t-1}$  we can discard  $f_{t-1}$  since we can always compute original predictions by  $h_{t-1}(g_{t \rightarrow t-1}(f_t(x)))$

# Online Generative Memory

Main idea:

- ▶ Idea is to keep model not to change its previous predictions
- ▶ Maybe we can adapt it to LwF scenario (but LwF does not work well even as it is)

Three ideas on how to achieve this

- ▶ MAML-like way:
  1. Perform  $k$  steps in  $\nabla \|f_{\theta}(x_k) - f_{\theta - \nabla L(\theta)}(x_k)\|_2^2$  direction
  2. Safely perform step in  $\nabla L(\theta)$  direction.
- ▶ Teacher distillation with previous batch after the current update
  1. Generate and save a batch of examples
  2. Perform gradient step for the main loss
  3. Perform teacher distillation step with the saved batch
  4. Repeat
- ▶ Project GD step onto  $\|f(x) - y\| \leq \varepsilon$  space (wrt spectral or frobenius norm)
  1. Generate and save a batch of examples  $X$
  2. Compute a gradient for the main loss
  3. Project the gradient onto  $f(x) = y$  loss by projecting the gradient for each layer