

1 Information Theory basics

Here we will give all the required definitions and their properties.

Definition 1.1. Entropy $H(X)$ of a discrete random variable X with probability mass function $p(X)$ is the quantity

$$H(X) = \mathbb{E}_{p(x)} [-\log p(x)]$$

Definition 1.2. Differential entropy of a random variable X with probability density function $p(x)$ is the quantity

$$H(X) = \mathbb{E}_{p(x)} [-\log p(x)]$$

2 Information Bottleneck as such

2.1 Origins and formulation

Information Bottleneck origins from rate-distortion theory, which is solving the following task: given a random variable X and some *distortion* function $d(x, \tilde{x})$ quantize X into \tilde{X} such that \tilde{X} is compressed as much as possible, but not too corrupted. Strictly speaking, we are trying to solve

$$\min_{p(\tilde{x}|x)} I(\tilde{X}; X) \quad \text{s.t.} \quad D(X, \tilde{X}) \leq D^*,$$

where

$$D(X, \tilde{X}) = \mathbb{E}_{p(x, \tilde{x})} [d(x, \tilde{x})]$$

and D^* is the maximum value of possible distortion that we permit. MSE loss or Hamming distance are common choices for $d(x, \tilde{x})$.

We can find optimal quantization by solving the variational problem:

$$\mathcal{L}(p(\tilde{x}|\tilde{x})) = I(X; \tilde{X}) + \beta D(X, \tilde{X})$$

The main problem with rate-distortion theory is that we need a distortion function $d(x, \tilde{x})$, which is difficult to specify for complex structured data, such as video or speech. And here Information Bottleneck comes to the rescue: what if after reconstruction we are only interested in some variable Y ? Then we can reformulate our problem as

$$\mathcal{L}(p(\tilde{x}|\tilde{x})) = I(X; \tilde{X}) - \beta I(Y, \tilde{X})$$

For discrete case it can be shown [1], that it's a special case of rate-distortion problem with $d(x, \tilde{x}) = D_{\text{KL}}[p(y|x) \parallel p(y|\tilde{x})]$.

2.2 A few words about sufficient statistics

- Two datasets which give us the same inference about sufficient statistic, would give the same inference about underlying parameter θ .
- Any injective function of sufficient statistic is also a sufficient statistic.
- Factorization theorem says, that $p(x|y) = h_T(X)g_T(T(X), y)$, dependence in g_T in x is only through $T(x)$.

2.3 Sufficiency and minimality of representations

From now on we'll use symbol Z for variable \tilde{X} , because it's more consistent with modern DL literature.

We have the following optimization problem:

$$\mathcal{L}(p(x; z)) = I(X; Z) - \beta I(Z; Y)$$

There is an interesting statement which connects notion of *sufficiency* and *minimality* of Z in the $\beta \rightarrow \infty$ setting. Let denote by $F(X)$ all random mappings of X (this means, that for $f \in F(X)$ we have Markov chain $Y \rightarrow X \rightarrow f(X)$). Let $S_X(Y)$ be a set of all sufficient statistics of X for Y .

Proposition 1. *If Z is a solution to*

$$\min_Z I(X; Z) \quad \text{s.t.} \quad I(Z; Y) = \max_{Z'} I(Z'; Y)$$

then Z is a minimal sufficient statistics of X for Y

$$P(X|Z, Y) = P(X|Z)$$

In other words, we are getting a minimal sufficient representation by optimizing Lagrangian in the limit $\beta \rightarrow \infty$. The proof of this theorem is split into two lemmas.

Lemma 1. *Z is a sufficient statistic of X for Y iff $I(Z; Y) = I(X; Y)$.*

Proof. Imagine that $T \in S_X(Y)$. For any $Z \in F(X)$ we have a Markov chain $Y \rightarrow X \rightarrow Z$, so by DPI we have $I(Y; Z) \leq I(Y; X)$. But by sufficient statistic property we have $P(X|Z, Y) = P(X|Y)$, so we have a Markov chain $Y \rightarrow Z \rightarrow X$. Again by DPI we have $I(Y; Z) \leq I(Y; X)$, so $I(Y; Z) = I(Y; X)$.

Now consider $Z = f(X)$ for some $f \in F(X)$ such that $I(Y; Z) = I(Y; X)$. As $Y \rightarrow X \rightarrow Z$ is a Markov chain, then by definition of conditional MI we have:

$$\begin{aligned} I(Y : Z|X) &\triangleq \mathbb{E}_X [\text{D}_{\text{KL}}[p(Y, Z|X) \parallel p(Y|X)p(Z|X)]] \\ &= \mathbb{E}_X [\text{D}_{\text{KL}}[p(Z|X, Y)p(Y|X) \parallel p(Y|X)p(Z|X)]] \\ &= \mathbb{E}_X [\text{D}_{\text{KL}}[p(Y|X)p(Z|X) \parallel p(Y|X)p(Z|X)]] = 0 \end{aligned}$$

Now, by chain rule for MI we have:

$$I(Y : X, Z) = I(Y : X) + I(Y, X|Z) = I(Y : X) + I(Y, Z|X) \implies I(Y, X|Z) = I(Y, Z|X) = 0$$

Applying definition of conditional MI again we get $p(Y, X|Z) = p(Y|Z)p(X|Z)$. And this means that $Z \in S_X(Y)$:

$$p(X|Y, Z) = \frac{p(X, Y|Z)}{p(Y|Z)} = \frac{p(X|Z)p(Y|Z)}{p(Y|Z)} = p(X|Z)$$

□

Let's denote by $S_X^*(Y)$ a set of minimal sufficient statistics of X for Y .

Lemma 2. *Let $Z \in S_X(Y)$, then*

$$Z \in S_X^*(Y) \iff I(X; Z) = \min_{T \in S_X(Y)} I(X; T)$$

Proof. First, let Z be a minimal sufficient statistic. Then for any other sufficient statistic T we have $Z = f(T)$ for some f . Then we get a Markov chain $X \rightarrow T \rightarrow Z$ and by DPI we have $I(X; Z) \leq I(X; T)$.

Now, let's prove reverse direction of the claim. Imagine, that $Z \in S_X(Y)$ but is not minimal. We are going to show, that $\exists T \in S_X(Y)$ such that $I(X; Z) > I(X; T)$. By Fisher-Neyman factorization theorem we have

$$Z \in S_X(Y) \iff \exists h_Z, g_Z \text{ s.t. } \forall x, y \quad p(x|y) = h_Z(x)g_Z(Z(x), y)$$

Let's define an equivalence relation

$$a \sim b \iff \forall y \exists \lambda \text{ s.t. } \frac{g_Z(a, y)}{g_Z(b, y)} = \lambda(a, b)$$

Now we define a deterministic function $T : \mathcal{X} \rightarrow \mathcal{Z}$ such that $\forall x : T(x) = \bar{z}$ — a representative of $[Z(x)]$ (TODO: exists by axiom of choice?). Let's prove, that it is a sufficient statistic. For this let's define

$$\begin{aligned} h_T(x) &\triangleq h_Z(x) \frac{g_Z(Z(x), y)}{g_Z(T(x), y)} \\ g_T(T(x), y) &\triangleq g_Z(T(x), y) \end{aligned}$$

Then we have

$$p(x|y) = h_Z(x)g_Z(Z(x), y) = h_Z(x)\frac{g_Z(Z(x), y)}{g_T(T(x), y)}g_T(T(x), y) = h_T(x)g_T(T(x), y)$$

Hence T is a sufficient statistic.

Now let's show that $I(X; Z) > I(X; T)$. Since Z is not minimal, then there is such $R \in S_X(Y)$ that Z is not a function of R . Let's show, that T is a function of R (btw, this will show, that T is minimal). For this we are going to show that if $R(x_1) = R(x_2)$ then $T(x_1) = T(x_2)$: this would allow us to build a function $\phi: \mathcal{R} \rightarrow \mathcal{T}$ which just take value r , find it's preimage $R^{-1}(r)$, take any sample $x \in R^{-1}(r)$ and compute $T(x)$. For any x_1, x_2 such that $R(x_1) = R(x_2)$ we have:

$$\begin{aligned} \frac{g_Z(Z(x_1), y)}{g_Z(Z(x_2), y)} &= \frac{p(x_1|y)h_Z(x_2)}{p(x_2|y)h_Z(x_1)} \\ &= \frac{h_R(x_1)g_R(R(x_1), y)h_Z(x_2)}{h_R(x_2)g_R(R(x_2), y)h_Z(x_1)} \\ &= \frac{h_R(x_1)g_R(R(x_1), y)h_Z(x_2)}{h_R(x_2)g_R(R(x_1), y)h_Z(x_1)} \\ &= \frac{h_R(x_1)h_Z(x_2)}{h_R(x_2)h_Z(x_1)} \\ &= \lambda(Z(x_1), Z(x_2)) \end{aligned}$$

This means that $Z(x_1) \sim Z(x_2)$, which in turn means that $T(x_1) = T(x_2)$, so T is minimal sufficient statistic and a function of R . \square

3 First steps: SZT experiments, critics

4 Disentangled representations

The following theorem gives us a hope to build invariant and disentangled representations. Informally, it says that if we can build sufficient and minimal representations (which we are building by optimizing IB Lagrangian, for example), then we get invariant and disentangled ones.

5 Information in the weights

5.1 Properties

5.2 Minimum description length

5.3 Connection to representations

6 Mutual information estimators

6.1 MINE

6.2 MI estimator under gaussian convolutions

7 Bibliography

References

- [1] Naftali Tishby, Fernando C. Pereira, and William Bialek. "The Information Bottleneck Method". In: 1999, pp. 368–377.

Exercises

1. Prove that

$$\arg \min_{q_1(x_1), \dots, q_n(x_n)} \text{D}_{\text{KL}}[p(x_1, \dots, x_n) \parallel q_1(x_1) \dots q_n(x_n)] = p(x_1) \dots p(x_n) \iff q_1(x_1), \dots, q_n(x_n) = p(x_1), \dots, p(x_n)$$

2. Prove that if X, Y are normally distributed, then Z is normally distributed too.
3. Maybe something on finding differential entropy of r.v. with infinite differential entropy. Or constructing such a variable.
4. Prove that $I(X; T)$ is either infinity or constant (except measure zero of weights).
5. Prove DPI
6. IB solution is a minimal sufficient statistic (based on <http://www.cs.huji.ac.il/labs/learning/Papers/ibgen.pdf>, theorem 5)