

# Ordered Neurons: Integrating Tree Structures into Recurrent Neural Networks<sup>1</sup>

May 21, 2020

---

<sup>1</sup> “Ordered Neurons: Integrating Tree Structures into Recurrent Neural Networks”  
by Shen et al.

# Overview

- ▶ Human language has a tree-like structure



- ▶ Human brain can infer it without any parse trees supervision
- ▶ How can we force neural networks to acquire this structure without supervision as well?
- ▶ Authors propose a specific inductive bias to do that:
  - ▶ divide a hidden state in a LSTM cell into chunks, each chunk keeps a hidden state for its own tree level
  - ▶ implement such a mechanism that can erase higher-level nodes only when all lower-level nodes are erased
  - ▶ this mechanism works by making forget/inputs gates of higher level nodes be dependent on the corresponding gates of lower-level ones
- ▶ They obtained good results for several tasks

# Vanilla LSTM

LSTM cell at timestep  $t$  takes an input  $x_t$ , updates its cell state  $c_t$  and hidden state  $h_t$  in the following way:

1. Compute the following quantities:

1.1  $f_t = \sigma(W_f x_t + U_f h_{t-1} + b_f)$  — “forget gate” to erase  $c_{t-1}$ ;

1.2  $i_t = \sigma(W_i x_t + U_i h_{t-1} + b_i)$  — “input gate” to erase new *raw* cell state  $\hat{c}_t$

1.3  $o_t = \sigma(W_o x_t + U_o h_{t-1} + b_o)$  — “output gate” to erase new hidden state  $h_t$

2. Then compute the following quantities

2.1  $\hat{c}_t = \tanh(W_c x_t + U_c h_{t-1} + b_c)$  — new “raw” cell state

2.2  $c_t = f_t \circ c_{t-1} + i_t \circ \hat{c}_t$  — new cell state

2.3  $h_t = o_t \circ \tanh(c_t)$  — new hidden state

Authors change the gating mechanism in such a way that erasement now occurs under some hierarchy.

# ON-LSTM (1/3)

- ▶ Imagine that we have some arbitrary tree and we observe only leaves (i.e. our model takes leaves as inputs one by one)
- ▶ At each timestep, we want to keep information not only about a leaf node, but also about every node on the path up to the root
- ▶ Let's store the information about each node of a path in a separate chunk of our state vector:

$$\mathbf{v} = [\underbrace{v_1^{(0)}, v_2^{(0)}, \dots, v_m^{(0)}}_{\text{leaf state}}, \underbrace{v_1^{(1)}, v_2^{(1)}, \dots, v_m^{(1)}}_{\text{leaf's parent state}}, \dots, \underbrace{v_1^{(L)}, v_2^{(L)}, \dots, v_m^{(L)}}_{\text{root state}}, ]$$

- ▶ We know that erasing is the core operation in LSTM to update its state
- ▶ So now we want a mechanism that would erase a higher level state only if all lower-level states are erased as well

## ON-LSTM (2/3)

- ▶ To implement the desired mechanism, our erasement vector  $\mathbf{g}$  for  $\mathbf{v}$  should have the following structure:

$$\mathbf{g} = [0, 0, 0, 0, \dots, 1, 1, 1],$$

i.e. we have a segment of zeros followed by a segment of ones.

- ▶ The idea is based on a novel activation function **cumax**:

$$\text{cumax}(\dots) = \text{cumsum}(\text{softmax}(\dots))$$

(it operates on vectors instead of scalars).

- ▶ Let  $d$  be a *split point*, i.e. a number when the first “1” occurred.
- ▶ Then  $d$  follows categorical distribution and we can model it with softmax:

$$p(d) = \text{softmax}(\hat{\mathbf{d}})$$

## ON-LSTM (3/4)

- ▶ Using  $d$  we can compute  $p(g_k = 1)$ , i.e. a probability of  $k$ -th element of  $g$  being equal to 1:

$$p(g_k = 1) = p(d \leq k) = \sum_{i \leq k} p(d = i)$$

- ▶ To compute a vector of probabilities  $[p(g_1 = 1), \dots, p(g_n = 1)]$  we need to apply cumax on  $\hat{d}$ :

$$p(\mathbf{g}) = \text{cumax}(\hat{\mathbf{d}})$$

- ▶ So we use a continuous relaxation for erasement operation (like in a vanilla LSTM) and erase a part of  $\mathbf{v}$  by computing

$$p(\mathbf{g}) \circ \mathbf{v}$$

- ▶ In this way, higher level segments are erased only if all the lower-level ones are erased

## ON-LSTM (4/5)

Imagine that at the current timestep we are at some level  $\ell$ , i.e.:

- ▶ we are going to “close” the branch of height  $\ell$
- ▶ we are going to erase the state below it
- ▶ we are going to keep all the higher-level context

# ON-LSTM (5/6)

First we should compute the master gates:

1.  $\tilde{f}_t = \text{cumax}(W_{\tilde{f}}x_t + U_{\tilde{f}}h_{t-1} + b_{\tilde{f}})$  — a mask, which discards all lower-level elements
2.  $\tilde{i}_t = 1 - \text{cumax}(W_{\tilde{i}}x_t + U_{\tilde{i}}h_{t-1} + b_{\tilde{i}})$  — a mask which keeps only lower-level elements
3. Compute the mask for the current-level segment  $\omega_t = \tilde{f}_t \circ \tilde{i}_t$ .

These vectors would look like this (ideally):

$$\begin{aligned}\tilde{f}_t &= [0, 0, 0, 0, 0, \textcolor{blue}{1}, \textcolor{blue}{1}, 1, 1] \\ \tilde{i}_t &= [1, 1, 1, 1, 1, \textcolor{blue}{1}, \textcolor{blue}{1}, 0, 0] \\ \omega_t &= [0, 0, 0, 0, 0, \textcolor{blue}{1}, \textcolor{blue}{1}, 0, 0]\end{aligned}\tag{1}$$



# ON-LSTM (6/6)

Then we update the states the following way:

3. Compute old element-wise forget and inputs gates  $f_t, i_t$  and new raw cell state  $\hat{c}_t$  without changes.
4.  $\hat{f}_t = f_t \circ \omega_t + (\tilde{f}_t - \omega_t)$  — forget gate that affects current level  $\ell$ , erases all lower-level elements and keeps all higher-level elements
5.  $\hat{i}_t = i_t \circ \omega_t + (\tilde{i}_t - \omega_t)$  — forget gate that affects current level  $\ell$ , erases all higher-level elements and keeps all lower-level elements
6.  $c_t = \hat{f}_t \circ c_{t-1} + \hat{i}_t \circ \hat{c}_t$  — new cell state
7.  $h_t$  is computed without changes

In this way we have updated only those part of the hidden state which corresponds to some specific level of a tree.

# Conclusion

- ▶ LSTM still performs better for short-term dependencies
- ▶ We repeat each dimension  $C$  times, before the element-wise multiplication
- ▶ It is interesting to see that there is room for exploration even in such an extensively explored model as LSTM
- ▶ Can we bring the idea into Transformers? Graph NNs?