

# ReZero is All You Need: Fast Convergence at Large Depth<sup>1</sup>

July 15, 2020

---

<sup>1</sup>*ReZero is All You Need: Fast Convergence at Large Depth* by Bachlechner et al., 2020

# Overview

- ▶ Initialization and stability is still an issue for many problems
- ▶ Authors propose a simple trick, similar to residual connections:

$$\mathbf{x}_{i+1} = \mathbf{x}_i + \alpha_i F(\mathbf{x}_i) \quad (1)$$

where  $\alpha_i$  is learnable and initialized at 0.

- ▶ It has the following benefits:
  - ▶ Simplicity and wide applicability
  - ▶ Faster convergence
  - ▶ It allows training of deeper models
- ▶ Authors test their approach on
  - ▶ Language modelling with Transformer
  - ▶ Classification on CIFAR-10
- ▶ They show good performance in terms of fast convergence and stability

# Residual with zero init (ReZero)

- Dynamical Isometry is a property that all singular values of the input-output Jacobian are close to 1
- It allows to train models much faster and make them much deeper
- Authors propose an easy trick that makes a model satisfy it (at init):

$$\mathbf{x}_{i+1} = \mathbf{x}_i + \alpha_i F(\mathbf{x}_i) \quad (2)$$

where  $\alpha_i$  is learnable and initialized at 0.




Figure 1: ReZero

- Experiments show that this property remains approximately preserved later on in training as well

# Fully-Connected models on CIFAR-10




Figure: Convergence speed for different normalization strategies

# Convolutional models on CIFAR-10

Model	Val. Error [%]	Change	Epochs to 80% Acc.	Train Loss $\times 1000$
ResNet-56 [2]	$6.27 \pm 0.06$	–	$20 \pm 1$	$5.9 \pm 0.1$
+ Gated ResNet [7, 29]	$6.80 \pm 0.09$	+ 0.53	$9 \pm 2$	$4.6 \pm 0.3$
+ zero $\gamma$ [23, 24]	$7.84 \pm 0.05$	+ 1.57	$39 \pm 4$	$31.2 \pm 0.5$
+ FixUp [10]	$7.26 \pm 0.10$	+ 0.99	$13 \pm 1$	$4.6 \pm 0.2$
+ <b>ReZero</b>	$6.58 \pm 0.07$	+ 0.31	$15 \pm 2$	$4.5 \pm 0.3$
ResNet-110 [2]	$6.24 \pm 0.29$	–	$23 \pm 4$	$4.0 \pm 0.1$
+ Gated ResNet [7, 29]	$6.71 \pm 0.05$	+ 0.47	$10 \pm 2$	$2.8 \pm 0.2$
+ zero $\gamma$ [23, 24]	$7.49 \pm 0.07$	+ 1.25	$36 \pm 5$	$18.5 \pm 0.9$
+ FixUp [10]	$7.10 \pm 0.22$	+ 0.86	$15 \pm 1$	$3.3 \pm 0.5$
+ <b>ReZero</b>	$5.93 \pm 0.12$	– 0.31	$14 \pm 1$	$2.6 \pm 0.1$
Pre-activation ResNet-18 [22]	$6.38 \pm 0.01$	–	$26 \pm 2$	$4.1 \pm 0.3$
+ <b>ReZero</b>	$5.43 \pm 0.06$	– 0.95	$12 \pm 1$	$1.9 \pm 0.3$
Pre-activation ResNet-50 [22]	$5.37 \pm 0.02$	–	$26 \pm 3$	$2.6 \pm 0.1$
+ <b>ReZero</b>	$4.80 \pm 0.08$	– 0.57	$17 \pm 1$	$2.2 \pm 0.1$

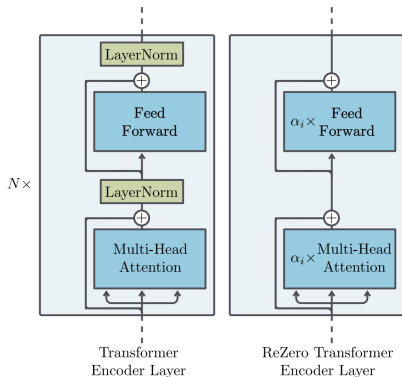
# ReZero Transformer

Vanilla Transformer uses Post-Norm normalization:

$$\mathbf{x}_{i+1} = \text{LayerNorm} (\mathbf{x}_i + \text{sublayer}(\mathbf{x}_i)) \quad (3)$$

Authors replaced this with:

$$\mathbf{x}_{i+1} = \mathbf{x}_i + \alpha_i \text{sublayer}(\mathbf{x}_i) \quad (4)$$



# Language Modeling results


Table 3: Comparison of various 12 layer Transformers normalization variants against ReZero and the training iterations required to reach 1.2 BPB on enwiki8 validation set.

Model	Iterations	Speedup
Post-Norm [27]	Diverged	-
+ Warm-up	13,690	$1\times$
Pre-Norm	17,765	$0.77\times$
GPT2-Norm [4]	21,187	$0.65\times$
ReZero $\alpha = 1$	14,506	$0.94\times$
<b>ReZero <math>\alpha = 0</math></b>	<b>8,800</b>	<b><math>1.56\times</math></b>

Table 4: Comparison of Transformers (TX) on the enwiki8 test set. Char-TX refers to the Character Transformer [14] and uses additional auxiliary losses to achieve its performance.

Model	Layers	Parameters	BPB
Char-TX [14]	12	41M	1.11
TX + Warm-up	12	38M	1.17
TX + ReZero $\alpha = 1$	12	34M	1.17
TX + ReZero $\alpha = 0$	12	34M	1.17
Char-TX [14]	64	219M	1.06
TX	64	51M	Diverged
TX + Warm-up	64	51M	Diverged
TX + ReZero $\alpha = 1$	64	51M	Diverged
TX + ReZero $\alpha = 0$	64	51M	1.11
TX + ReZero	128	101M	1.08

# Model preserves dynamic isometry by itself



**Figure:** Histograms of  $\log(\sigma)$  of singular values. Left: traditional Transformer. Right: 64-layer ReZero Transformer



## Residual weights evolution




Figure: Evolution of  $\alpha_i$  for 64-layer Transformer

- ▶ Model first increases  $\alpha_i$  for later layers, then decreases them all.
- ▶ Authors say that there is a similar pattern for  $\alpha = 1$  (for a 12-layer transformer): model first tries to reduce  $\alpha$ . But instead of increasing